

# Final Review

---

Max Turgeon

STAT 3150–Statistical Computing

# Main theme

- Recall the main theme of the course: **using computational techniques to solve statistical problems.**
- What kind of statistical problems?
  - Point estimation
  - Interval estimation
  - Hypothesis testing

# Numerical methods and Optimisation

- For the first two modules, we specifically looked at point estimation.
- We talked about the following methods:
  - Bisection/Brent's method for root finding in one dimension.
  - Newton-Raphson for optimisation in any dimension.

# Generating random variates

- R has many built-in functions for generating random variates.
  - `runif`, `rnorm`, etc.
- We discussed general techniques when these functions aren't enough.
  - Inverse transform, or generally any type of transformation.
  - Accept-reject sampling.
- **When would you need to generate random variates?**
  - Estimate expected values (i.e. Monte Carlo integration)
  - Estimate probability statements
  - Simulation studies

# Accept-Reject algorithm i


- Recall the general goal: we want to sample from a distribution with density  $f$ .
- Instead, we sample from a “dominating” density  $g$ , and then randomly decide if we keep the sample or not.
  - **Dominating:** There exists a constant  $c > 1$  such that  $\frac{f(t)}{g(t)} \leq c$  for all  $t$  such that  $f(t) > 0$ .
  - In particular, the support of  $g$  must include that of  $f$  (it’s best if they are equal).
- To decide if we keep the sample, we sample  $u$  from a uniform  $U(0, 1)$  and accept if  $u < \frac{f(y)}{cg(y)}$ .
- Suppose we want to sample from a complicated distribution:

## Accept-Reject algorithm ii

```
dfn <- function(x){  
  out <- numeric(length(x))  
  out[x <= 0.5] <- dnorm(x[x <= 0.5])  
  out[x > 0.5] <- dnorm(x[x > 0.5]) +  
    dchisq(x[x > 0.5], df = 2)  
  out  
}
```

```
# Let's visualize  
x_vec <- seq(-5, 5, length.out = 1000)  
plot(x_vec, dfn(x_vec), type = 'l')
```

## Accept-Reject algorithm iii



# Accept-Reject algorithm iv


- The support is the whole real line. This gives us some natural candidates for the dominating distribution:
  - Normal distribution
  - t distribution
- Goal: Find a dominating (or proposal) distribution with “heavier tails” so that the ratio  $f(t)/g(t)$  is bounded.
- Let's try the normal density.

```
# Visualize the ratio
```

```
plot(x_vec, dfn(x_vec)/dnorm(x_vec), type = 'l')
```



# Accept-Reject algorithm v




## Accept-Reject algorithm vi

- The normal density will not work, because the ratio is unbounded...
- Let's try the Cauchy distribution (aka  $t(1)$ ):

```
# Visualize the ratio
```

```
plot(x_vec, dfn(x_vec)/dcauchy(x_vec), type = 'l')
```

## Accept-Reject algorithm vii



## Accept-Reject algorithm viii

- It looks like the ratio is bounded above by  $c = 4$ !
- We can double-check using calculus: if the ratio is bounded, then it will attain a maximum, which corresponds to when the derivative is equal to 0.
- The density of interest is piece-wise defined: we would technically need to optimize on both pieces, but from our figure we know the maximum will be on the second piece (e.g.  $x > 0.5$ ).
  - We will omit the details. You can either solve with Calculus, or use a numerical method to find a more exact  $c$ .

## Accept-Reject algorithm ix


```
# Set parameters----  
C <- 4 # Constant  
n <- 1000 # Number of variates  
k <- 0 # counter for accepted  
j <- 0 # iterations  
y <- numeric(n) # Allocate memory
```

## Accept-Reject algorithm $x$

```
while (k < n) {  
  u <- runif(1)  
  j <- j + 1  
  x <- rcauchy(1) # random variate from g  
  if (u < dfn(x)/(C*dcauchy(x))) {  
    k <- k + 1  
    y[k] <- x  
  }  
}
```

```
# Visualize the histogram  
hist(y, 50, freq = FALSE)
```

# Accept-Reject algorithm xii





- **Note:** We didn't even start with an actual density; it doesn't integrate to 1.
- But the algorithm still works (proof left as an exercise).

# Monte Carlo integration i

- This topic mostly falls under *point estimation*.
- Estimate quantities of the form

$$E(g(X)) = \int g(x)f(x)dx, \quad X \sim f.$$

- Trace plot = diagnose convergence issues
- Variance reduction
  - Antithetic variables
  - Control variates
  - Importance sampling
- Confidence intervals in MC integration are based on the Central Limit Theorem

- Since our estimates are sample means, we need to divide by  $\sqrt{n}$ , where  $n$  is the number of variates in the sample mean.
- **When would you use MC integration?**
  - To estimate difficult integrals.
  - Many, many estimators can be defined as expected values of transformations  $g(X)$  of a random variable  $X$ .

# Importance sampling

- It's a form of **variance reduction** for Monte Carlo integration.
- Based on the following identity:

$$E_f(g(X)) = E_\phi \left( \frac{g(X)f(X)}{\phi(X)} \right),$$

as long as  $\phi$  is nonzero on the support of  $f$ .

- We want to choose the importance function  $\phi$  such that:
  - $\phi$  is a density from which it is “easy” to sample.
  - the ratio  $\frac{|g(X)f(X)|}{\phi(X)}$  is almost constant.
- **Why do we care so much about reducing variance anyway?**
  - Because smaller variance means smaller confidence intervals, which means more accurate inference.

# How to integrate?

Where does MC integration fit within the different ways of doing integration?

Choose the first thing that works in this order:

1. Exact integration (works well for simple integrand)
2. Symbolic mathematics (e.g. Maple or Mathematica)
3. Numerical integration (works best in 1 dimension)
4. **MC Integration**
5. Markov Chain Monte Carlo (cf. Bayesian statistics)

# Importance Sampling—Example i

- Consider the following integral:

$$\int_0^\pi \frac{dx}{x^2 + \cos^2 x}.$$

- How can we write this as an expected value  $E_f(g(X))$ ?

## Importance Sampling—Example ii

- We can take:
  - $g(x) = 1/(x^2 + \cos^2 x)$
  - $X \sim \text{Unif}(0, \pi)$

```
g_fun <- function(x) pi/(x^2 + cos(x)^2)
f_vars <- runif(1000, min = 0, max = pi)
mean(g_fun(f_vars))
```

```
## [1] 1.542862
```

## Importance Sampling—Example iii

- In importance sampling, we now need an importance function  $\phi(x)$ , which is the density of a distribution we can sample from.
- Let's try  $\phi$  the density of  $Exp(1)$ .
- **Important observation:** The density  $f$  of  $Unif(0, \pi)$  is zero outside the integral bounds, so our ratio  $g(x)f(x)/\phi(x)$  will be zero as expected. If it wasn't, we would need to modify our implementation of  $g$  to make sure it is.



## Importance Sampling—Example iv

```
phi_vars <- rexp(1000)
# Think of re-weighting
weights <- dunif(phi_vars,
                  max = pi)/dexp(phi_vars)
mean(g_fun(phi_vars)*weights)

## [1] 1.584207
```

## Importance Sampling—Example v

- You could improve the performance of this estimator by sampling from a *truncated* Exponential.
- You could use an Exponential with a different mean (perhaps also truncated).
- These two are left as an exercise.

# Monte Carlo methods for Inference

- This module was an interlude, connecting Monte Carlo integration and resampling methods.
  - What is a statistic? An estimator? A sampling distribution?
  - What is a type I error? Type II error? Power?
- If we are willing to completely specify the data generating mechanism, we can study the consequences of these assumptions through **Monte Carlo simulation**.
  - Which estimator is more efficient (i.e. has smallest variance)?
  - Does my confidence interval have the right coverage probability?
  - Which hypothesis test has largest power?

# Resampling methods

- The next few modules were on resampling methods:
  - *Jackknife*: “Resample” all subsets of size  $n - 1$ .
  - *Bootstrap*: Resample  $n$  observations **with replacement**.
  - *Permutation tests*: Permute all observations to mimic resampling under the null hypothesis
- It’s during these modules that we finally started analysing data.

# Jackknife

- Mainly presented for its historical importance.
  - It can be formalized as a “linear approximation” to the bootstrap.
  - It also helps motivate some quantities/techniques, e.g. student residuals, Cook’s distance, leave-one-out cross-validation.
- We used it for estimating the **standard error** and **bias** of an estimate.
  - Use the formulas provided.
  - But it doesn’t always work! E.g. median, quantiles.
- We can construct confidence intervals using the CLT.
  - No general accuracy guarantees.

- Bootstrap is almost always preferable to jackknife.
  - It is valid under more general assumptions.
  - Can be used to construct valid confidence intervals.
- Recall the general idea: we are trying to mimic going back and collecting more samples.
  - **What should we bootstrap?** The quantities we sampled in the first place.
  - In practice, we need to resample rows of `data.frames`, so that we *preserve* the correlation structure between the different measurements.
- We discussed 5 types of confidence intervals for bootstrap.

## Bootstrap ii

- **Remember:** no need to divide by  $\sqrt{n}$  with bootstrap.
- **Very important:** Throughout, we assumed simple random sampling.
  - In particular, we assumed the observations were independent.
  - Can be generalized with care.
- We spent considerable time talking about bootstrap and linear regression.
  - *Resample residuals:* Need to assume linearity, additivity, equal variance AND independence of the errors.
  - *Resample cases/rows:* Only need to assume independence.

- If the only thing you remember from STAT 3150 after the final exam is how to *properly* use bootstrap, I'll be happy.
  - But it won't be enough to pass the final exam!



# Permutation tests

- Jackknife and bootstrap are mainly used for point/interval estimation.
  - Of course, any confidence interval leads to a hypothesis test.
- Permutation tests are specifically for hypothesis testing.
- They are a family of **non-parametric** methods for testing

$$H_0 : F = G.$$

- In other words: two-sample tests of equal distribution.
- **Main idea:** if the data all come from the same distribution, then which ones are Xs and which ones are Ys is irrelevant.

- Statistical computing is deeply connected to modern statistics.
  - You **cannot** do statistics in the 21st century without computing.
- Some areas of statistics are more computational than others:
  - Statistical learning
  - Bayesian statistics
  - High-dimensional data